


Chapter 3: Processes





Chapter 3: Processes


- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication (IPC)
 - Examples of IPC Systems
 - Communication in Client-Server Systems





Objectives

- To introduce the notion of a process -- a program in execution, which forms the basis of all computation
- To describe the various features of processes, including scheduling, creation and termination, and communication
- To explore interprocess communication using shared memory and message passing
- To describe communication in client-server systems





Process Concept

- An operating system executes a variety of programs:
 - Batch system – **jobs**
 - Time-shared systems – **user programs** or **tasks**
- Textbook uses the terms ***job*** and ***process*** almost interchangeably
- Program is ***passive*** entity stored on disk (**executable file**), process is ***active***
 - Program becomes process when executable file loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc






Process Concept

- One program can be several processes
 - Consider multiple users executing the same program
- **Process** – a program in execution; process execution must progress in sequential fashion
- Process includes multiple parts
 - The program code, also called **text section**
 - Current activity involving **program counter**, processor registers
 - **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time





Process in Memory






Process State

- As a process executes, it changes **state**
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution





Diagram of Process State





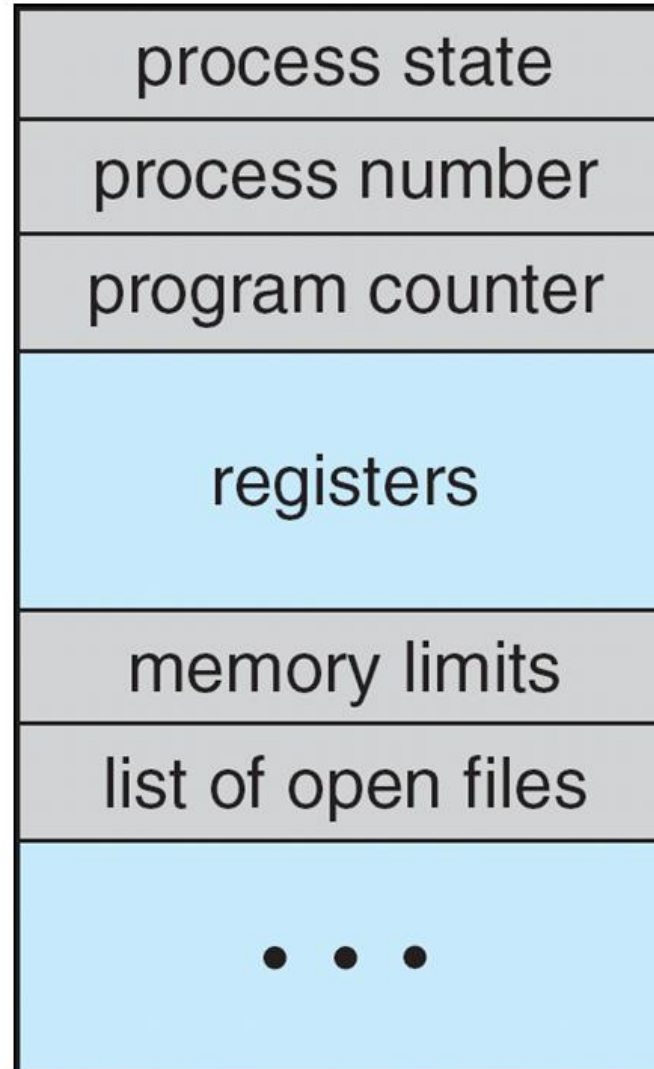
Process Control Block (PCB)

- Information associated with each process
 - **Process state** – running, waiting, etc
 - **Program counter** – location of instruction to next execute
 - **CPU registers** – contents of all process-centric registers
 - **CPU scheduling information** – priorities, scheduling queue pointers
 - **Memory-management information** – memory allocated to the process
 - **Accounting information** – CPU used, clock time elapsed since start, time limits
 - **I/O status information** – I/O devices allocated to process, list of open files







Process Control Block (PCB)





CPU Switch From Process to Process





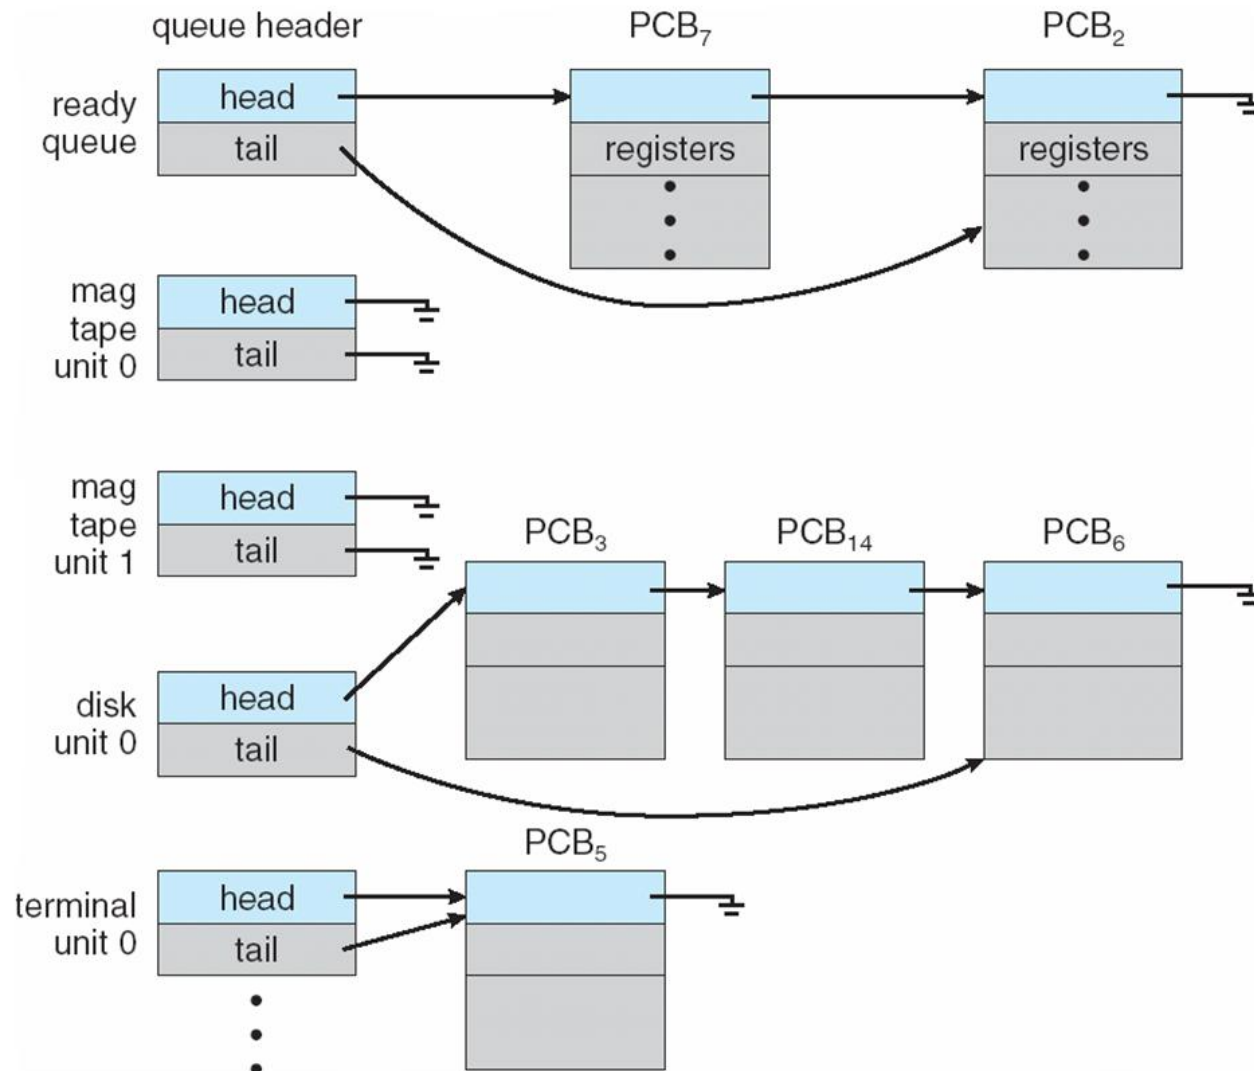
Process Scheduling

- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler** selects among available processes for next execution on CPU
- Maintains **scheduling queues** of processes
 - **Job queue** – set of all processes in the system
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - **Device queues** – set of processes waiting for an I/O device
- Processes migrate among the various queues





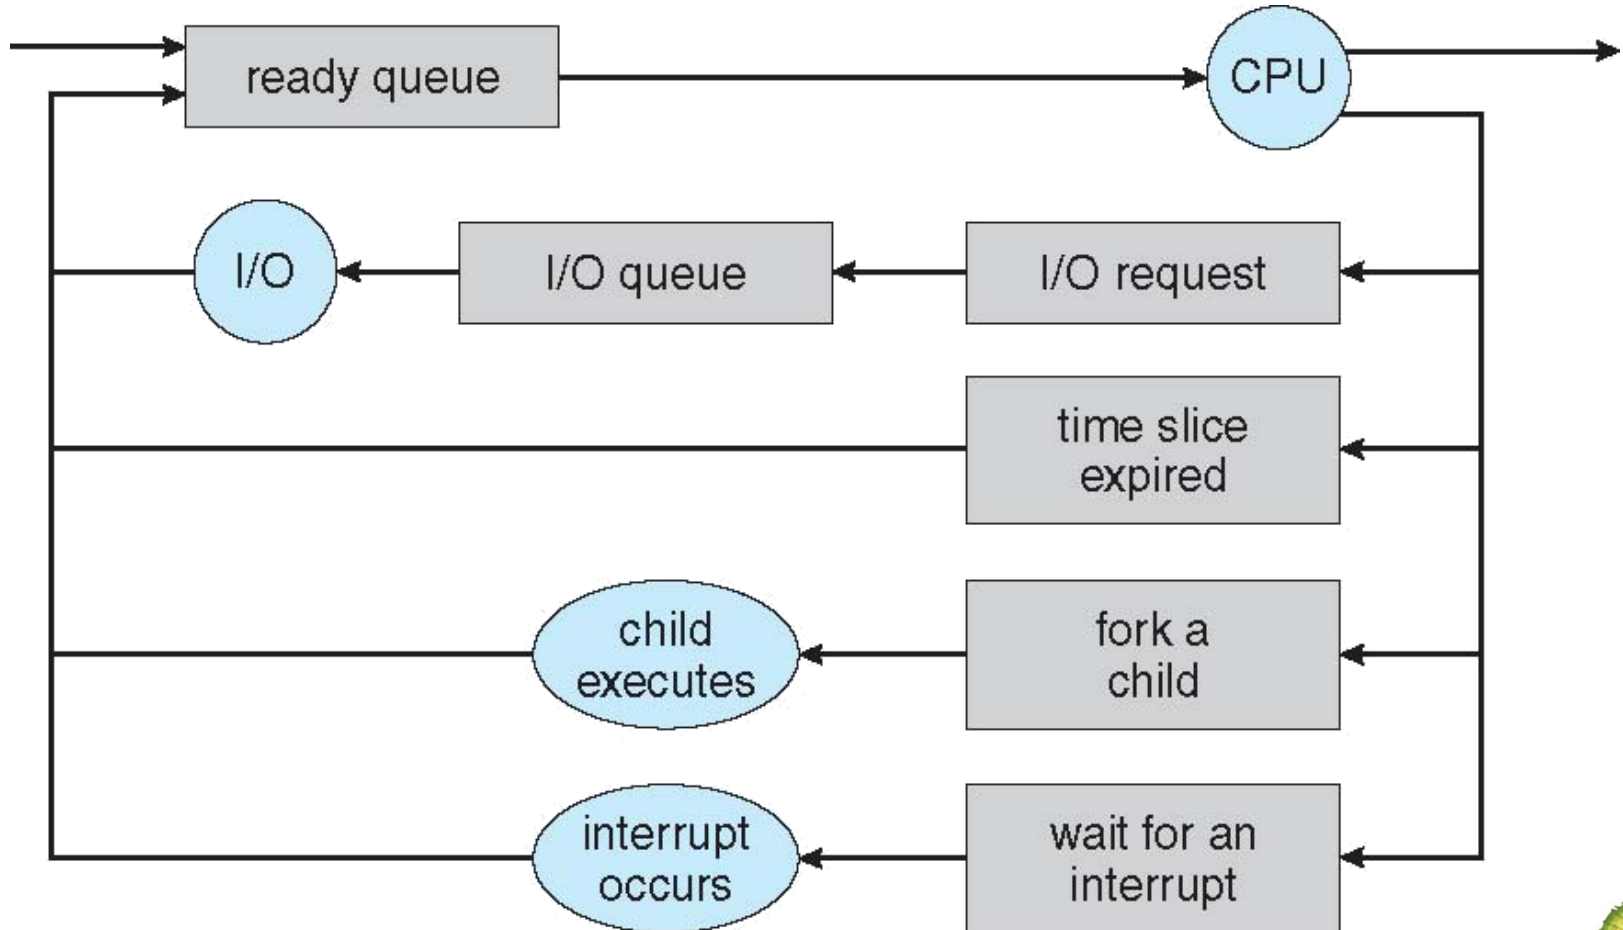
Ready Queue and Various I/O Device Queues





Representation of Process Scheduling

Queuing diagram represents queues, resources, flows





Schedulers

- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
- **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
- The long-term scheduler controls the **degree of multiprogramming**





Schedulers

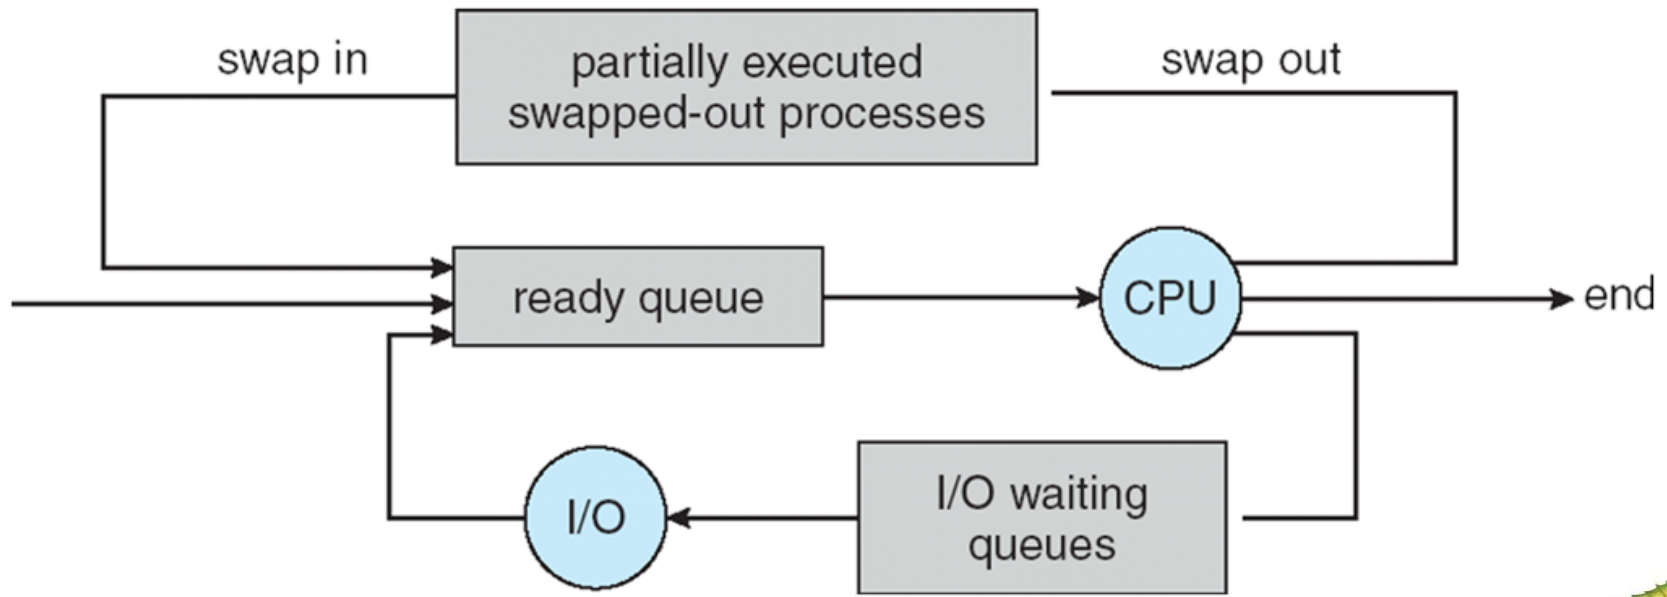
- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast)
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- Long-term scheduler strives for good **process mix**





Addition of Medium Term Scheduling

- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
 - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**





Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
 - Context of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB
 - => longer the context switch
- Time dependent on hardware support





Operations on Processes

- System must provide mechanisms for process creation, termination, and so on as detailed next






Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
 - Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until all or some children terminate





A Tree of Processes in Linux






Process Creation (Cont.)

■ Address space

- Child duplicate of parent program
- Child has a program loaded into it

■ UNIX examples

- **fork()** system call creates new process
- **exec()** system call used after a **fork()** to replace the process' memory space with a new program





C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```





Process Termination

- Process executes last statement and asks the operating system to delete it (**exit()**)
 - Child process return its status to parent (via **wait()**)
 - Process' resources are deallocated by the operating system
- Parent may terminate execution of its children processes (**abort()**) for a variety of reasons:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - ▶ Some operating systems do not allow child to continue if its parent terminates
 - ▶ All children must also terminated - **cascading termination**





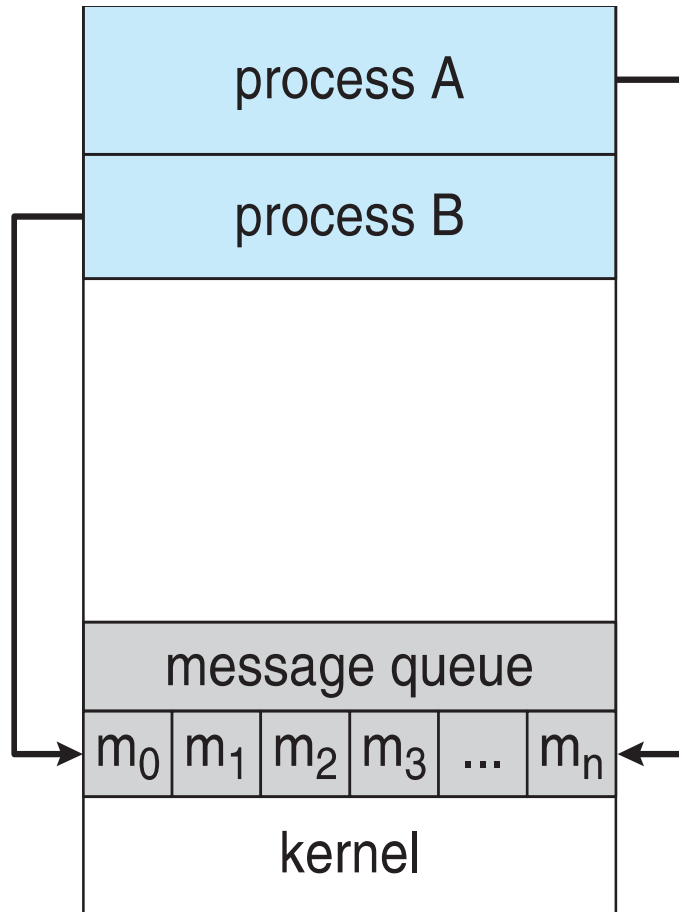
Interprocess Communication

- Processes within a system may be **independent** or **cooperating**
- Cooperating process can affect or be affected by other processes, including sharing data
 - Reasons for cooperating processes:
 - ▶ Information sharing
 - ▶ Computation speedup
 - ▶ Modularity
 - ▶ Convenience
- Cooperating processes require an **interprocess communication (IPC)** mechanism
 - Two models of IPC:
 - a) **Message passing**
 - b) **Shared memory**






Communications Models



(a)



(b)





Producer-Consumer Problem

- Paradigm for cooperating processes, a **producer** process produces information that is consumed by a **consumer** process:
 - A compiler may produce assembly code that is consumed by an assembler
 - A web server produces HTML files which are consumed by the client web browser
- One solution to allow them to run concurrently, uses a buffer of items (**shared memory**) that can be filled by producer and emptied by consumer:
 - **unbounded-buffer** places no practical limit on the size of the buffer
 - **bounded-buffer** assumes that there is a fixed buffer size





Bounded-Buffer – Shared-Memory Solution

■ Shared data

```
#define BUFFER_SIZE 10  
typedef struct {  
    . . .  
} item;  
  
item buffer[BUFFER_SIZE];  
int in = 0;  
int out = 0;
```

■ Solution is correct, but can only use BUFFER_SIZE-1 elements





Bounded-Buffer – Producer

```
item next_produced;

while (true) {

    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing -- no free buffers */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```





Bounded-Buffer – Consumer

```
item next_consumed;

while (true) {

    while (in == out)

        ;    /* do nothing - empty buffers */

    next_consumed = buffer[out];

    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next_consumed */

}
```





Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides at least two operations:
 - **send**(*message*) – message size fixed or variable
 - **receive**(*message*)
- If P and Q wish to communicate, they need to:
 - establish a **communication link** between them
 - exchange messages via send/receive
- Implementation of communication link
 - logical (e.g., direct or indirect, synchronous or asynchronous, automatic or explicit buffering)





Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?





Direct Communication

- Processes must name each other explicitly:
 - **send** (P , *message*) – send a message to process P
 - **receive**(Q , *message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional





Indirect Communication

- Messages are directed and received from **mailboxes** (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional





Indirect Communication

■ Operations

- Create a new mailbox
 - ▶ A mailbox may be created (and then owned) either by a process or by the operating system
- Send and receive messages through mailbox
 - ▶ Only owner process can receive messages through its mailbox
- Destroy a mailbox by the owner

■ Primitives are defined as:

- **send**(*A, message*) – send a message to mailbox A
- **receive**(*A, message*) – receive a message from mailbox A





Indirect Communication

■ Mailbox sharing

- P_1 , P_2 , and P_3 share mailbox A owned by OS
- P_1 sends; P_2 and P_3 receive
- Who gets the message?

■ Solutions


- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was





Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue
 - **Non-blocking receive** has the receiver receive a valid message or null





Synchronization (Cont.)

- Different combinations possible
 - If both send and receive are blocking, we have a **rendezvous**
- Producer-consumer problem becomes trivial

```
Producer  message next_produced;
           while (true) {
               /* produce an item in next_produced */
               send(next_produced);
           }
```

```
Consumer message next_consumed;
           while (true) {
               receive(next_consumed);
               /* consume the item in next_consumed */
           }
```





Buffering

- Queue of messages attached to the link; implemented in one of three ways:

1. Zero capacity – 0 messages

Sender must wait for receiver (rendezvous)

2. Bounded capacity – finite length of n messages

Sender must wait if link full

3. Unbounded capacity – infinite length

Sender never waits





Communications in Client-Server Systems

- Shared memory and message passing strategies can be used for communication in client–server systems as well
- Four other used techniques for communication in client–server systems:
 - Sockets
 - Remote Procedure Calls
 - Pipes
 - Remote Method Invocation (Java)





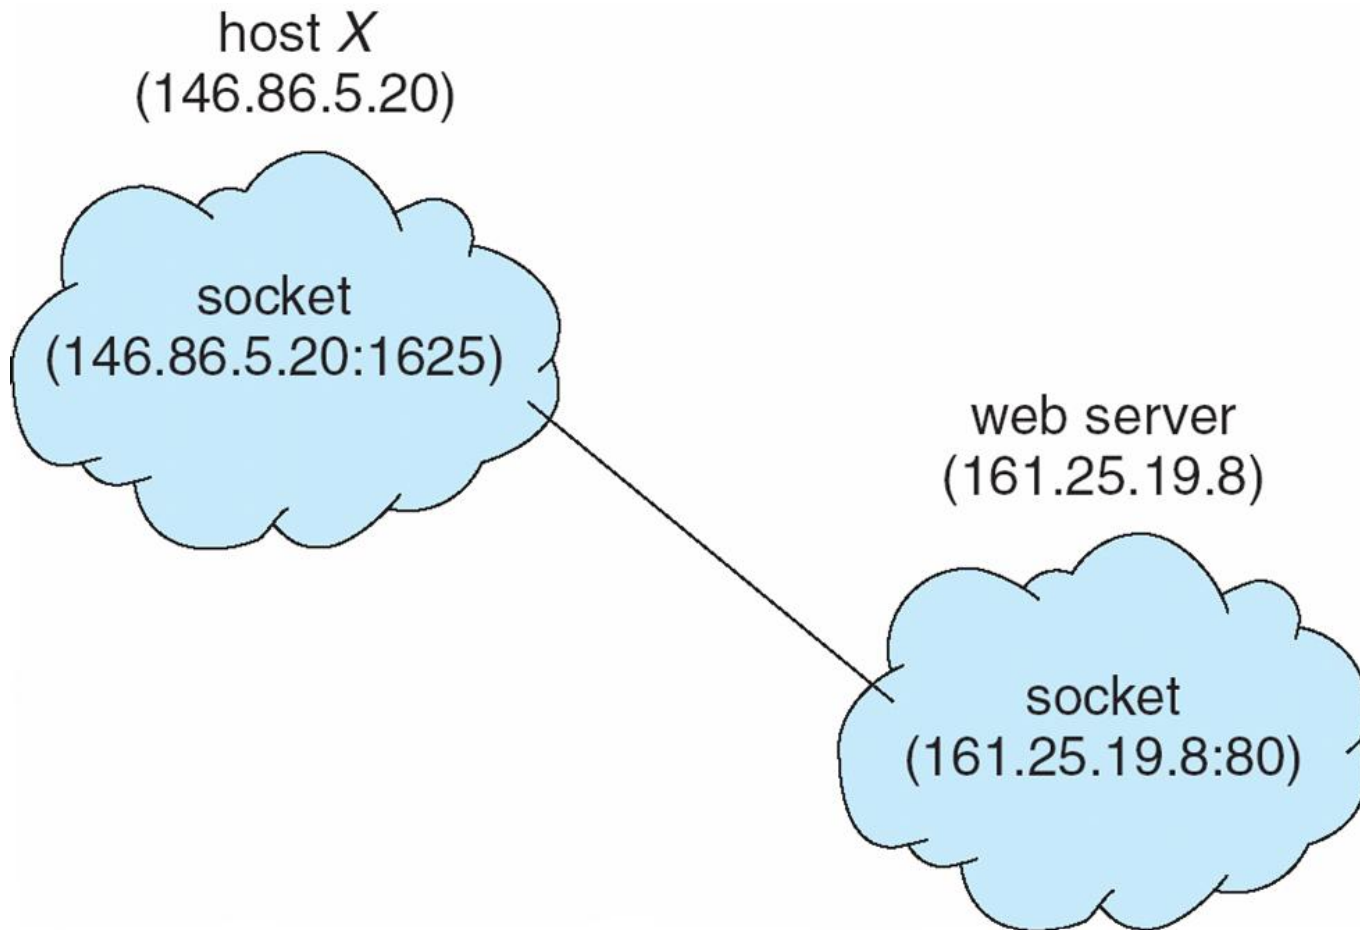
Sockets

- A **socket** is an endpoint for communication
 - A pair of processes communicating over a network employs a pair of sockets
- A socket is identified by an IP address concatenated with a **port** number.
 - The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- The server waits for incoming client requests by listening to a specified port
 - All ports below 1024 are **well known**, used for standard services
 - To allow a client and server on the same host to communicate, a special IP address 127.0.0.1 (**loopback**) is used to refer to itself





Socket Communication





Sockets in Java

- Three different types of sockets:
 - **Connection-oriented (TCP) Socket** class
 - **Connectionless (UDP) DatagramSocket** class
 - **MulticastSocket** class – data can be sent to multiple recipients
- Consider this “Date” server:

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





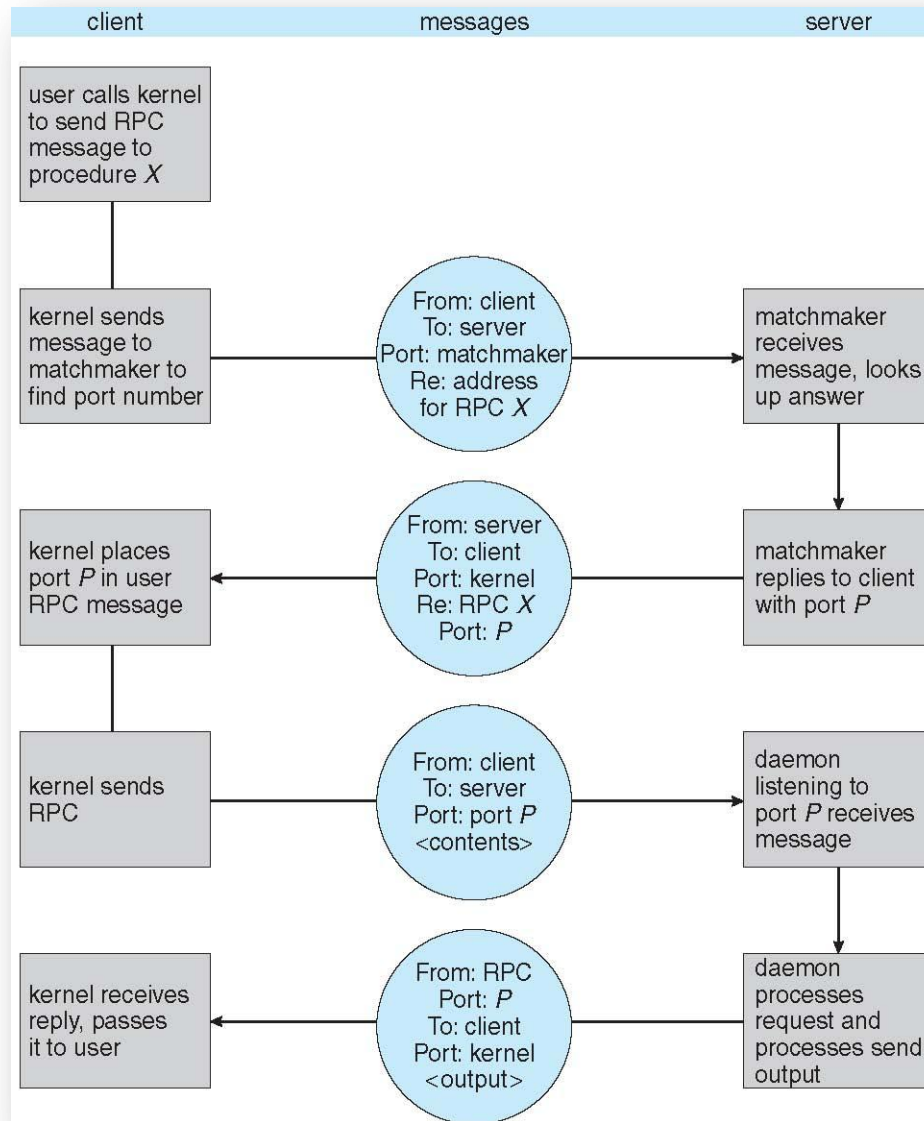
Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
 - Again uses ports for service differentiation
- **Stubs** – client-side proxy for the actual procedure on the server
 - The client-side stub locates the server and **marshalls** the parameters
 - The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server





Execution of RPC






Pipes

- Acts as a channel allowing two processes to communicate
- Simpler ways for communication but have some limitations
- **Issues**
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (as **parent-child**) between the communicating processes?
 - Can be used over a network or must reside on the same machine?






Ordinary Pipes



- Ordinary Pipes allow communication in standard producer-consumer style
 - Producer writes to one end (the **write-end** of the pipe)
 - Consumer reads from the other end (the **read-end** of the pipe)
 - So, they allow only unidirectional communication
- Require parent-child relationship between communicating processes
 - A parent process creates a pipe and uses it to communicate with its child process
- UNIX treats it as a special type of file and can be accessed using `read()` and `write()` system calls





Named Pipes


- More powerful tool than ordinary pipes
 - Communication can be bidirectional
 - No parent–child relationship is required.
 - Several processes can use it for communication
- Both UNIX and Windows systems support it, although implementation details differ greatly
 - Referred to as FIFOs in UNIX and once created, they appear as typical files and manipulated with ordinary `open()`, `read()`, `write()`, and `close()` system calls
 - Windows provide a richer communication mechanism
 - ▶ Full-duplex communication is allowed
 - ▶ May reside on either the same or different machines
 - ▶ Allow either byte- or message-oriented data






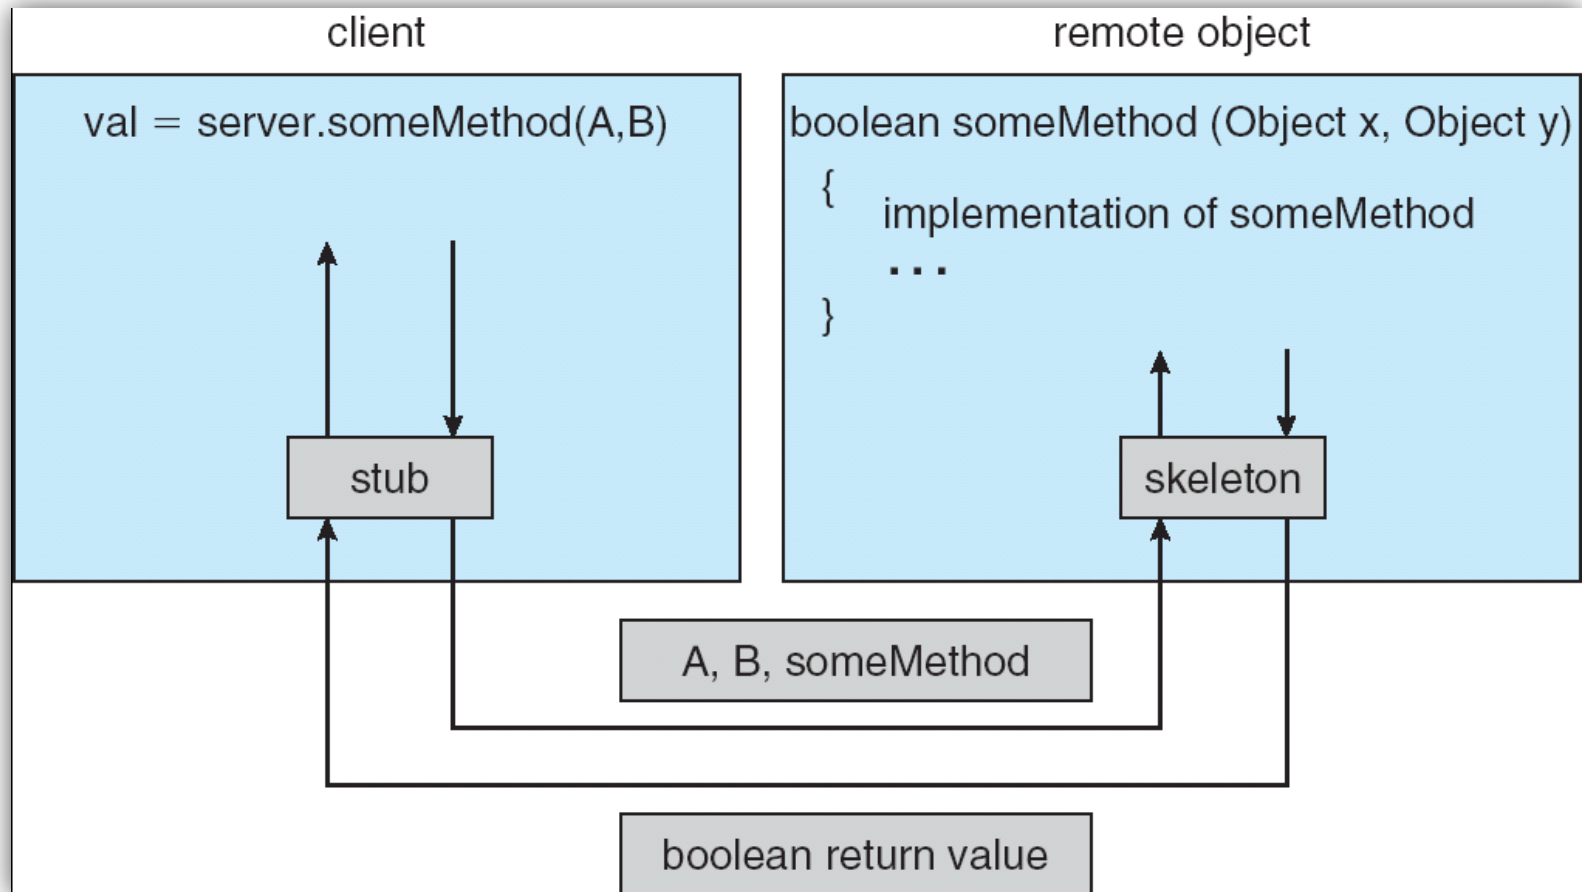
Remote Method Invocation

- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs.
- RMI allows a Java program on one machine to invoke a method on a remote object.





Marshalling Parameters



End of Chapter 3

