


# Chapter 8: Main Memory

---





# Chapter 8: Memory Management

---

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table
- Examples:
  - The Intel 32 and 64-bit Architectures
  - ARM Architecture





# Objectives

---

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation
- To provide a detailed description of the Intel Pentium, which supports both pure segmentation and segmentation with paging





# Background

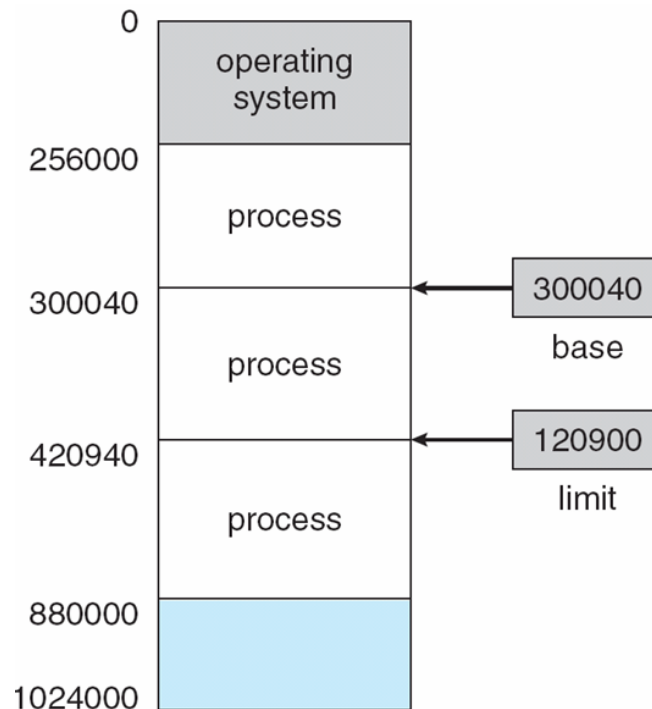
- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are only storage CPU can access directly
  - Memory unit only sees a stream of addresses + read requests, or address + data and write requests
  - Register access in one CPU clock (or less)
  - Main memory can take many cycles, causing a **stall**
  - **Cache** sits between main memory and CPU registers
- Protection of memory required to ensure correct operation






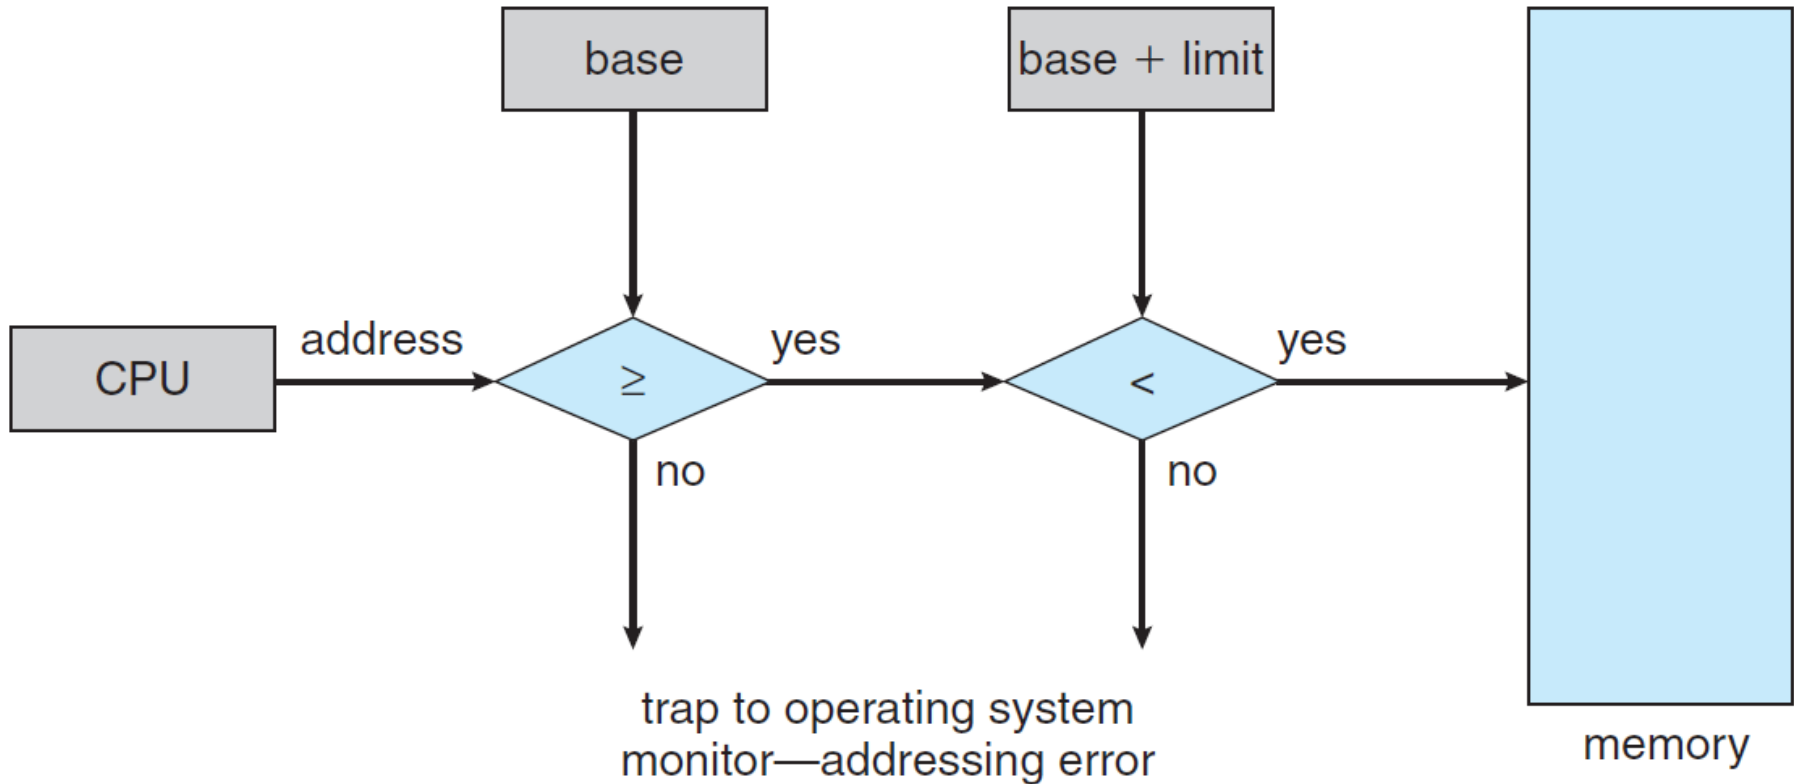
# Base and Limit Registers

- A pair of **base** (relocation) and **limit registers** define the logical address space
- CPU must check, every memory access generated in user mode, to be sure it is between base and limit for that user





# Hardware Address Protection with Base and Limit Registers





# Address Binding

- Programs on disk, ready to be brought into memory to execute form an **input queue**
  - Without support, must be loaded into address 0000
- Inconvenient to have first user process physical address always at 0000
  - How can it not be?
- Further, addresses represented in different ways at different stages of a program's life
  - Source code addresses usually symbolic
  - Compiled code addresses **bind** to relocatable addresses (i.e. "14 bytes from beginning of this module")
  - Linker or loader will bind relocatable addresses to absolute addresses (i.e. 74014)
  - Each binding maps one address space to another





# Binding of Instructions and Data to Memory

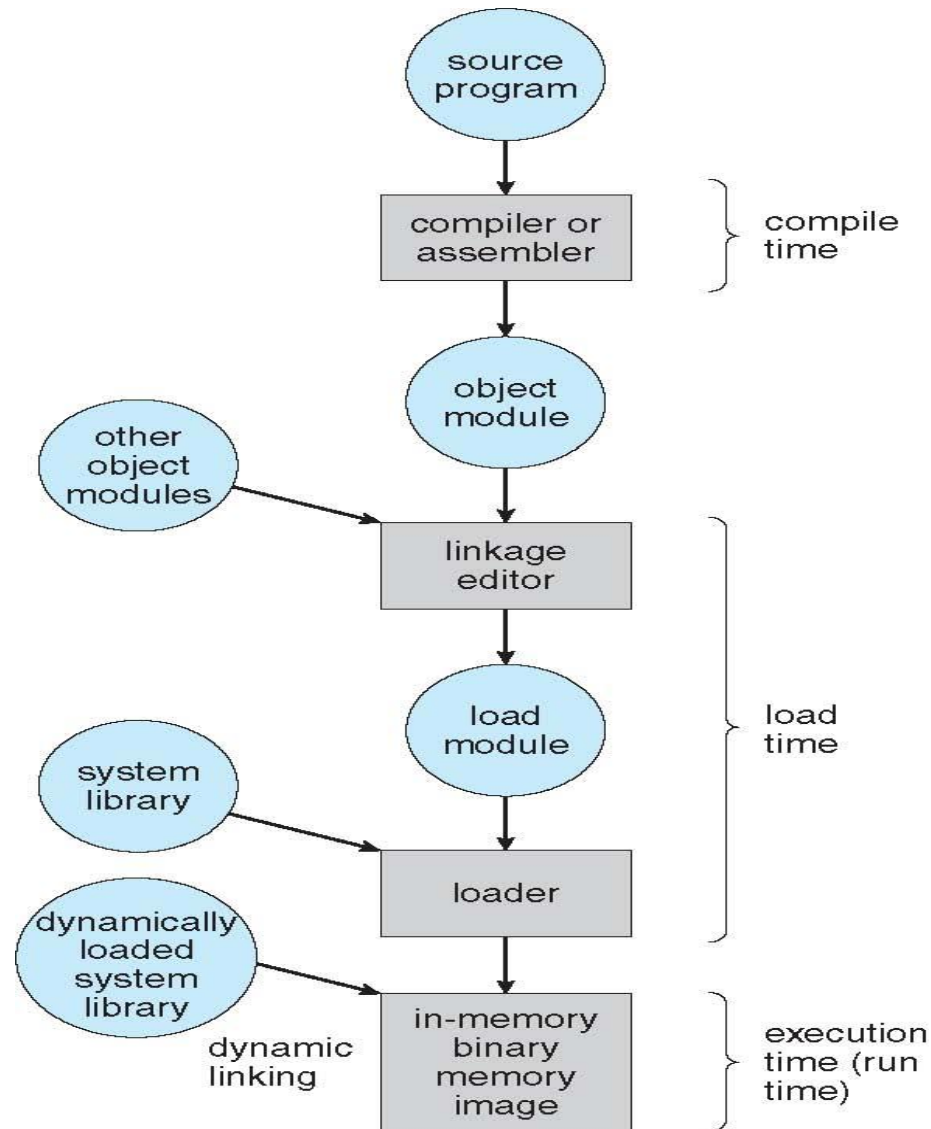
- Address binding of instructions and data to memory addresses can happen at three different stages:
  - **Compile time:** If memory location known a priori, **absolute code** can be generated
    - ▶ Must recompile code if starting location changes
  - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
  - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
    - ▶ Need hardware support for address maps (e.g., base and limit registers)







# Multistep Processing of a User Program





# Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
  - **Logical address** – generated by the CPU; also referred to as **virtual address**
  - **Physical address** – address seen by memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
  - **Logical address space** is the set of all logical addresses generated by a program
  - **Physical address space** is the set of all physical addresses generated by a program






# Memory-Management Unit (MMU)

- Hardware device that at run time maps logical (virtual) to physical address
  - Many methods possible, covered in rest of this chapter
- To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
  - Base register now called **relocation register**
  - MS-DOS on Intel 80x86 used 4 relocation registers
- User program deals with *logical* addresses; it never sees the *real* physical addresses
  - Execution-time binding occurs when reference is made to location in memory







# Dynamic Relocation Using a Relocation Register

- Routine is not loaded until it is called
  - Better memory-space utilization; unused routine is never loaded
  - All routines kept on disk in relocatable load format
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
  - Implemented through program design
  - OS can help by providing libraries to implement **dynamic loading**





# Dynamic Relocation Using a Relocation Register





# Dynamic Linking

- **Static linking** – system libraries combined by the loader into the binary program image
- **Dynamic linking** – linking of system libraries postponed until execution time
  - Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
  - Stub replaces itself with the address of the routine, and executes the routine
  - Operating system checks if routine is in processes' memory address
    - ▶ If not in address space, add to address space
- **Dynamic linking is particularly useful for libraries**
  - This system is also known as **shared libraries**
  - This feature can be extended to library updates (such as bug fixes) that may be replaced by a new version






# Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
  - Total physical memory space of processes can exceed physical memory
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed





# Schematic View of Swapping







# Swapping

- Major part of swap time is transfer time
  - total transfer time is directly proportional to the amount of memory swapped
- System maintains a **ready queue** consisting of all ready-to-run processes whose memory images are on backing store or in memory
- Does the swapped out process need to swap back in to same physical addresses?
  - Depends on address binding method
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
  - Swapping normally disabled and started only if more than threshold amount of memory allocated
  - Then disabled again once memory demand reduced below threshold





# Context Switch Time Including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
  - Context switch time can then be very high
  - 100 MB process swapping to hard disk with transfer rate of 50 MB/sec
    - ▶ Swap out time of 2000 ms
    - ▶ Plus swap in of same sized process
    - ▶ Total context switch swapping component time of 4000 ms
  - Can reduce time through decreasing size of memory swapped – by knowing how much memory really being used
    - ▶ System calls to inform OS size of used memory via `request_memory()` and `release_memory()`





# Swapping

- Other constraints as well on swapping
  - Pending I/O – can't swap out as I/O would occur to wrong process
  - Or transfer I/O to kernel space, then to I/O device
    - ▶ Known as **double buffering**, but adds overhead
- Standard swapping not used in modern OSs
  - But modified version common
  - Swap only when free memory extremely low
- Swapping not supported on mobile systems
  - Instead use other methods to free memory if low space
    - ▶ Android terminates apps if low free memory
    - ▶ iOS **asks** apps to voluntarily release allocated memory





# Contiguous Allocation

---

- Main memory must accommodate both OS and various user processes
  - Must to allocate main memory in most efficient way
  - Contiguous allocation is one early method
- Main memory usually divided into two **partitions**:
  - Resident operating system, usually held in low memory with interrupt vector
  - User processes then held in high memory
  - Each process contained in a single contiguous section of memory






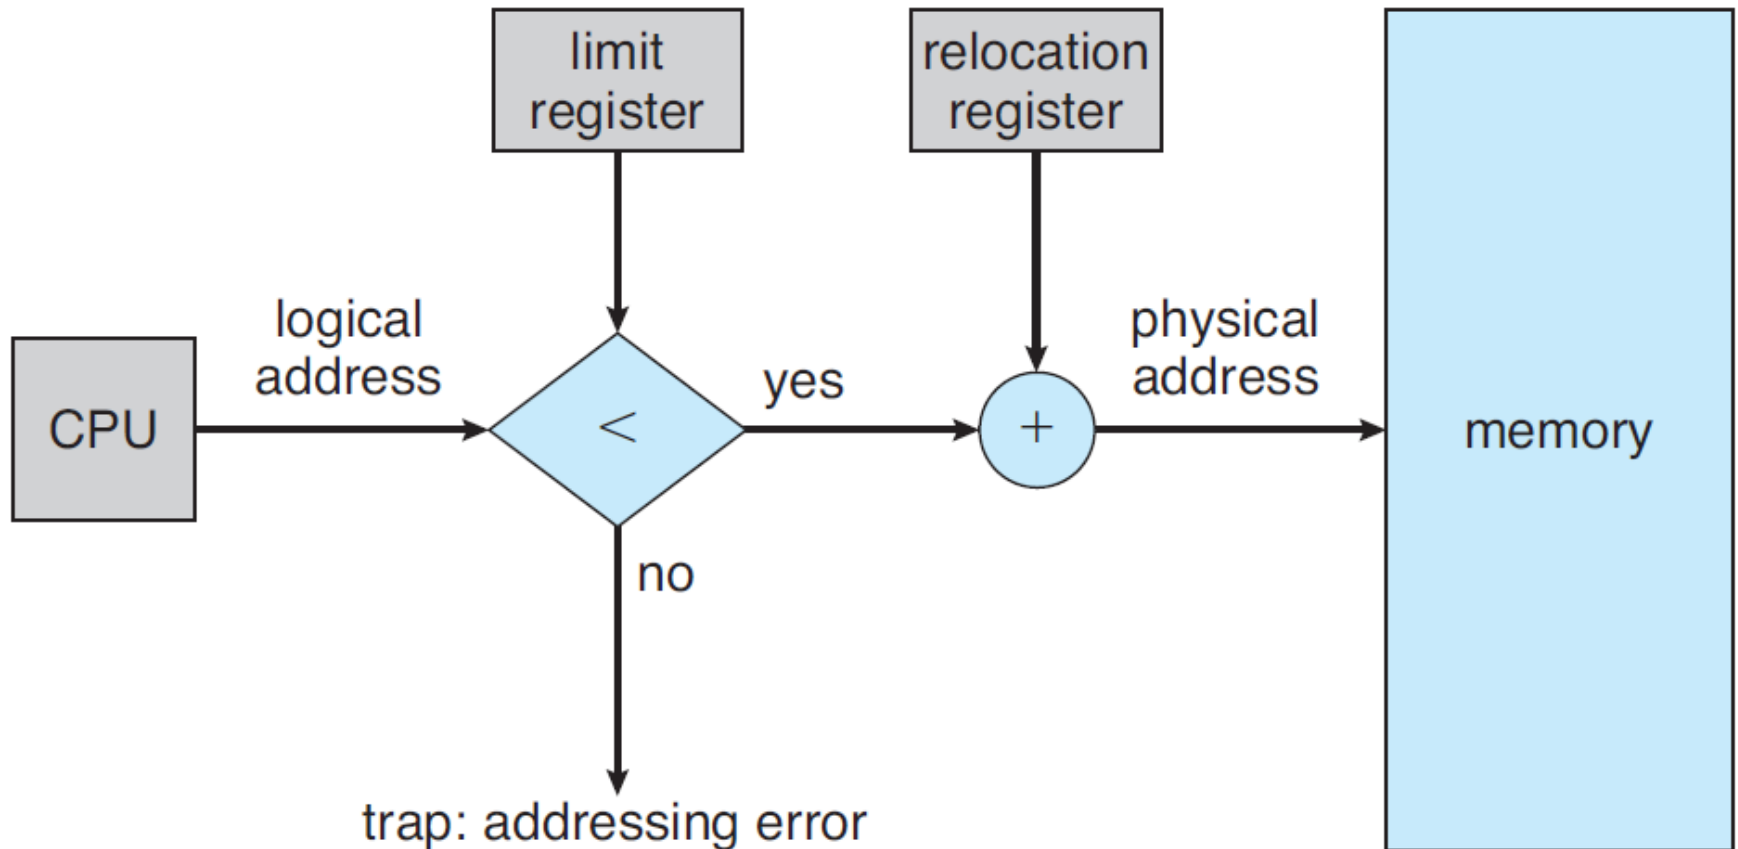
# Contiguous Allocation

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
  - Base register contains value of the smallest physical address
  - Limit register contains range of logical addresses
    - ▶ Each logical address must be less than the limit register
  - MMU maps logical address *dynamically* by adding the value in the base register
  - As part of context switch, dispatcher loads base and limit registers values of a process under execution
- This scheme provides an effective way to allow operating system's size to change dynamically.
  - Suitable if an OS service is not commonly used
  - Such services code is called **transient** code





# Hardware Support for Relocation and Limit Registers





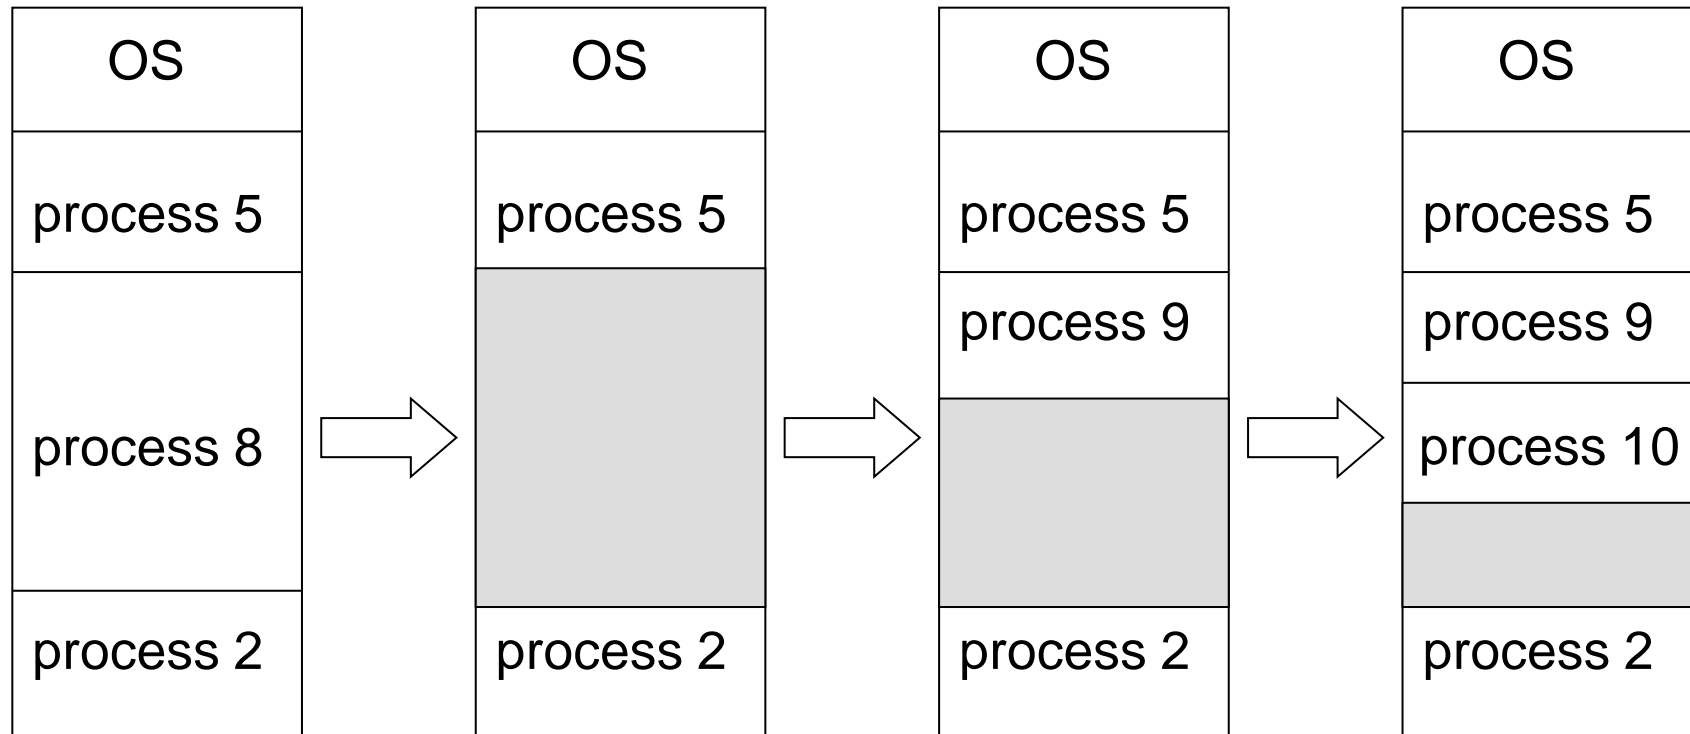
# Contiguous Allocation

- Memory is divide into several partitions – **multiple-partition method**
  - Degree of multiprogramming limited by number of partitions
  - **Variable-partition** sizes for efficiency (sized to a given process' needs)
  - **Hole** – block of available memory; holes of various size are scattered throughout memory
  - When a process arrives, it is allocated memory from a hole large enough to accommodate it
  - Process exiting frees its partition and adjacent free partitions combined
  - Operating system maintains information about:
    - a) allocated partitions
    - b) free partitions (hole)





# Contiguous Allocation







# Dynamic Storage-Allocation Problem

How to satisfy a request of size  $n$  from a list of free holes?

- **First-fit**: Allocate the *first* hole that is big enough
- **Best-fit**: Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
  - Produces the smallest leftover hole
- **Worst-fit**: Allocate the *largest* hole; must also search entire list
  - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization





# Fragmentation

- **External Fragmentation** – total memory space is enough to satisfy a request, but it is not contiguous
  - Storage is fragmented into a large number of small holes
  - First-fit and best-fit strategies for memory allocation suffer from external fragmentation
  - First fit analysis reveals that given  $N$  blocks allocated,  $1/2 N$  blocks lost to fragmentation
    - ▶  $1/3$  of memory may be unusable  $\Rightarrow$  **50-percent rule**
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used





# Fragmentation (Cont.)

- Reduce external fragmentation by **compaction**
  - Shuffle memory contents to place all free memory together in one large block
  - Compaction is possible *only* if relocation is dynamic, and is done at execution time
    - ▶ Relocation requires just moving the process and then changing the base register to reflect the new base address
  - I/O problem
    - ▶ Latch job in memory while it is involved in I/O
    - ▶ Do I/O only into OS buffers
- Now consider that backing store has same fragmentation problems





# Segmentation


---


- Memory-management scheme that mapped the programmer's view of memory
- A program is a collection of segments and each segment is a logical unit, such as:
  - main program
  - procedure
  - function
  - method
  - object
  - local variables, global variables
  - common block
  - stack
  - symbol table
  - arrays






# Programmer's View of a Program






# Logical View of Segmentation



Logical address space



Physical memory space






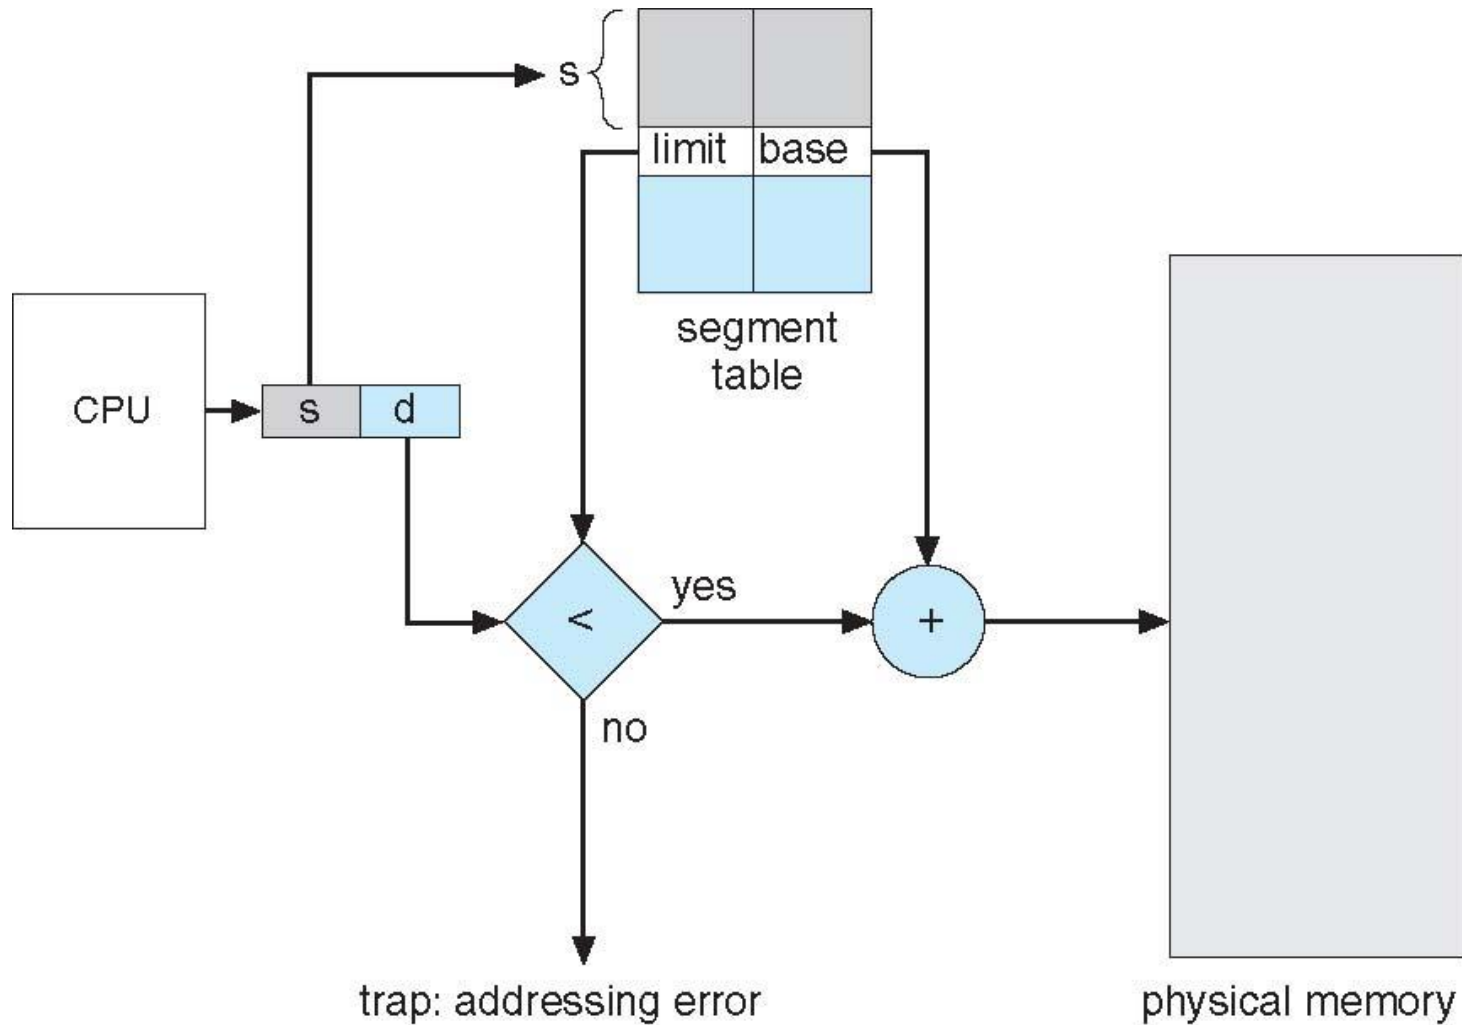
# Segmentation Architecture

- Logical address consists of a two tuple:  
    <segment-number, offset>
- A **Segment table** is used to map the two-dimensional logical addresses of each process into one-dimensional physical addresses
  - Each entry in the segment table has:
    - ▶ **base** – contains the starting physical address where the segment reside in memory
    - ▶ **limit** – specifies the length of the segment
- Hardware support:
  - **Segment-table base register (STBR)** points to the segment table's location in memory
  - **Segment-table length register (STLR)** indicates number of segments used by a process;
    - ▶ segment number **s** is legal if **s** < **STLR**





# Segmentation Hardware







# Segmentation Architecture

## ■ Protection

- With each entry in segment table associate:

- ▶ validation bit = 0  $\Rightarrow$  illegal segment
- ▶ read/write/execute privileges

## ■ Protection bits associated with segments; code sharing occurs at segment level


## ■ Since segments vary in length, memory allocation is a dynamic storage-allocation problem

## ■ A segmentation example is shown in the following diagram






# Example of Segmentation



	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

segment table





# Paging

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
  - Avoids external fragmentation
  - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
  - Size is power of 2
  - Normally, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **pages**





# Paging

---


- Keep track of all free frames
- To execute a program of size ***N*** pages, need to find ***N*** free frames and load program
- Set up a **page table** to translate logical addresses to physical addresses
  - Map the corresponding physical frame for each logical page
- Backing store likewise split into pages
- Still have Internal fragmentation
  - Last frame allocated may not be completely full





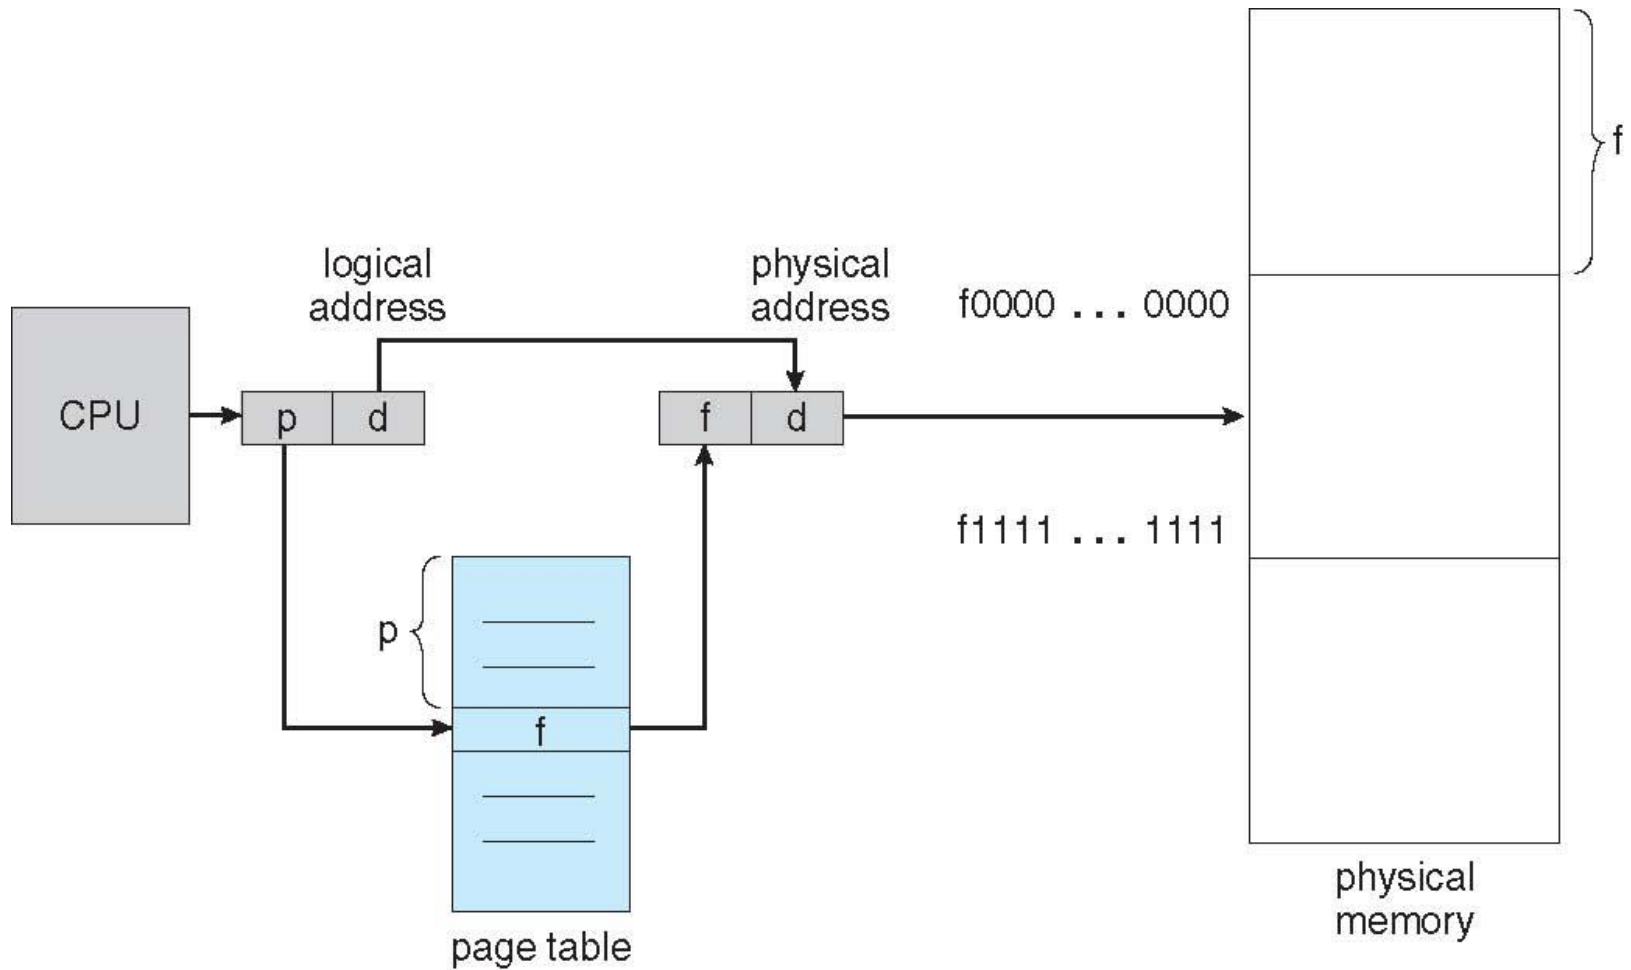
# Address Translation Scheme

- Address generated by CPU is divided into:
  - **Page number** ( $p$ ) – used as an index into a page table which contains base address of each page in physical memory
  - **Page offset** ( $d$ ) – combined with base address to define the physical memory address that is sent to the memory unit
- For given logical address space  $2^m$  and page size  $2^n$ :






# Paging Hardware





# Paging Model of Logical and Physical Memory




logical  
memory

0	1
1	4
2	3
3	7

page table

frame  
number



physical  
memory





# Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

A 32-byte memory with 4-byte pages ( $n=2$  and  $m=4$ )








# Paging

- Process view and physical memory now very different
- By implementation, process can only access its own memory
- Calculating internal fragmentation:
  - Page size = 2,048 bytes
  - Process size = 72,766 bytes
  - 35 pages + 1,086 bytes
  - Internal fragmentation of  $2,048 - 1,086 = 962$  bytes
  - Worst case fragmentation = 1 frame – 1 byte
  - On average fragmentation =  $1 / 2$  frame size
  - So small frame sizes desirable?
    - ▶ But each page table entry takes memory to track
    - ▶ Page sizes growing over time






# Free Frames



Before allocation


After allocation





# Implementation of Page Table

- Dedicated registers are used for implementing the page table if the table is reasonably small
  - In DEC PDP-11, the address consists of 16 bits, and the page size is 8 KB.
  - The page table thus consists of eight entries that are kept in fast registers.
- Most computers allow the page table to be very large; therefore, the page table is kept in main memory
  - **Page-table base register (PTBR)** points to the page table
  - **Page-table length register (PTLR)** indicates size of the page table





# Implementation of Page Table

- In the last scheme every data/instruction access requires two memory accesses
  - One for the page table and one for the data/instruction
- The two memory access problem can be solved using a special fast lookup memory cache called **translation look-aside buffer (TLB)**
  - Each entry in TLB consists of two parts: a page no. as a key and a corresponding frame value
  - An item is compared with all keys simultaneously
  - If a page number is not in the TLB (a **TLB miss**), its frame is added for faster access next time
    - ▶ Replacement policies must be considered





# Associative Memory

## ■ Associative memory – parallel search

Page #	Frame #

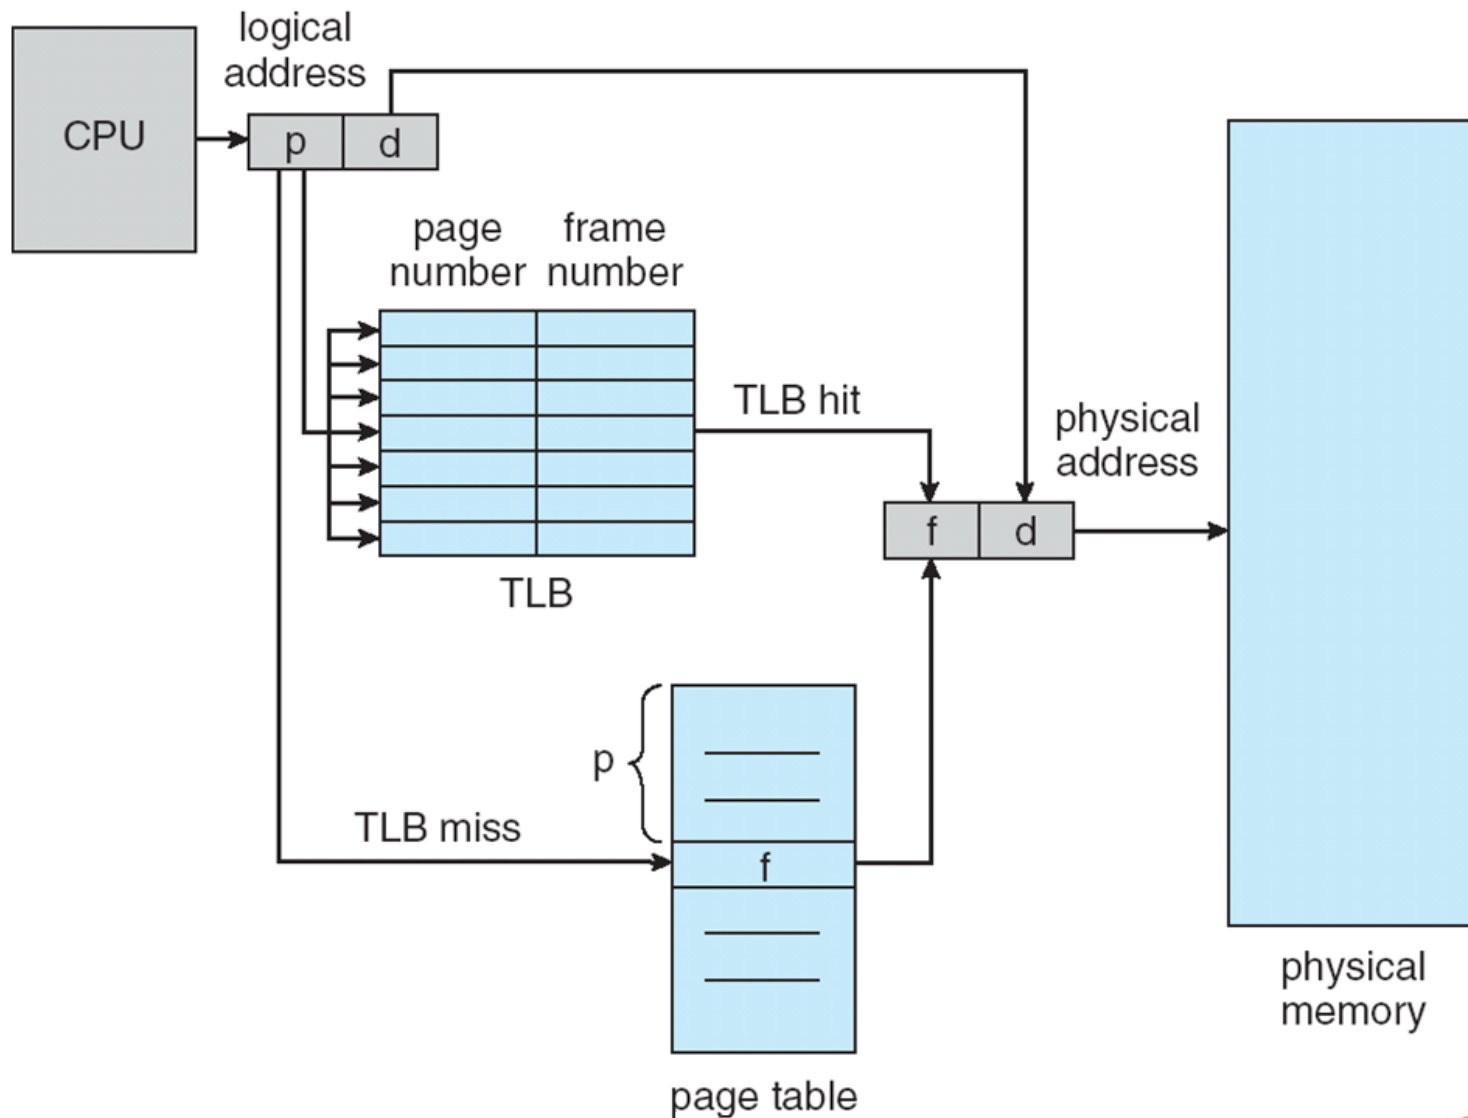
## ■ Address translation (p, d)

- If p is in associative register, get frame # out
- Otherwise get frame # from page table in memory





# Paging Hardware With TLB





# Effective Access Time

- Associative Lookup =  $\varepsilon$  time unit
  - Can be  $< 10\%$  of memory access time
- Hit ratio =  $\alpha$ 
  - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers

- **Effective Access Time (EAT)**

$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$


- Consider  $\alpha = 80\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - ▶  $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio  $\alpha = 99\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - ▶  $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$





# Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
  - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and it is a legal page
  - “invalid” indicates that the page is not in the process’ logical address space
  - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel







# Valid (v) or Invalid (i) Bit In A Page Table

00000	page 0
	page 1
	page 2
	page 3
	page 4
10,468	page 5
12,287	

frame number		valid-invalid bit
0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

page table

0	
1	
2	page 0
3	page 1
4	page 2
5	
6	
7	page 3
8	page 4
9	page 5
	⋮
	page <i>n</i>





# Shared Pages

---

## ■ Shared code

- One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed

## ■ Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space





# Shared Pages Example

ed 1
ed 2
ed 3
data 1

process  $P_1$

3
4
6
1

page table  
for  $P_1$

ed 1
ed 2
ed 3
data 2

process  $P_2$

3
4
6
7

page table  
for  $P_2$


ed 1
ed 2
ed 3
data 3

process  $P_3$

3
4
6
2

page table  
for  $P_3$

0	
1	data 1
2	data 3
3	ed 1
4	ed 2
5	
6	ed 3
7	data 2
8	
9	
10	
11	





# Structure of the Page Table


- Memory structures for paging can get huge using straight-forward methods
  - Consider a 32-bit logical address space as on modern computers
  - Page size of 4 KB ( $2^{12}$ )
  - Page table would have 1 million entries ( $2^{32} / 2^{12}$ )
  - If each entry is 4 bytes  $\Rightarrow$  4 MB of physical address memory space for page table alone
    - ▶ That amount of memory used to cost a lot
    - ▶ Don't want to allocate that contiguously in main memory
- Alternative methods:
  - Hierarchical Paging
  - Hashed Page Tables
  - Inverted Page Tables





# Hierarchical Page Tables



- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table





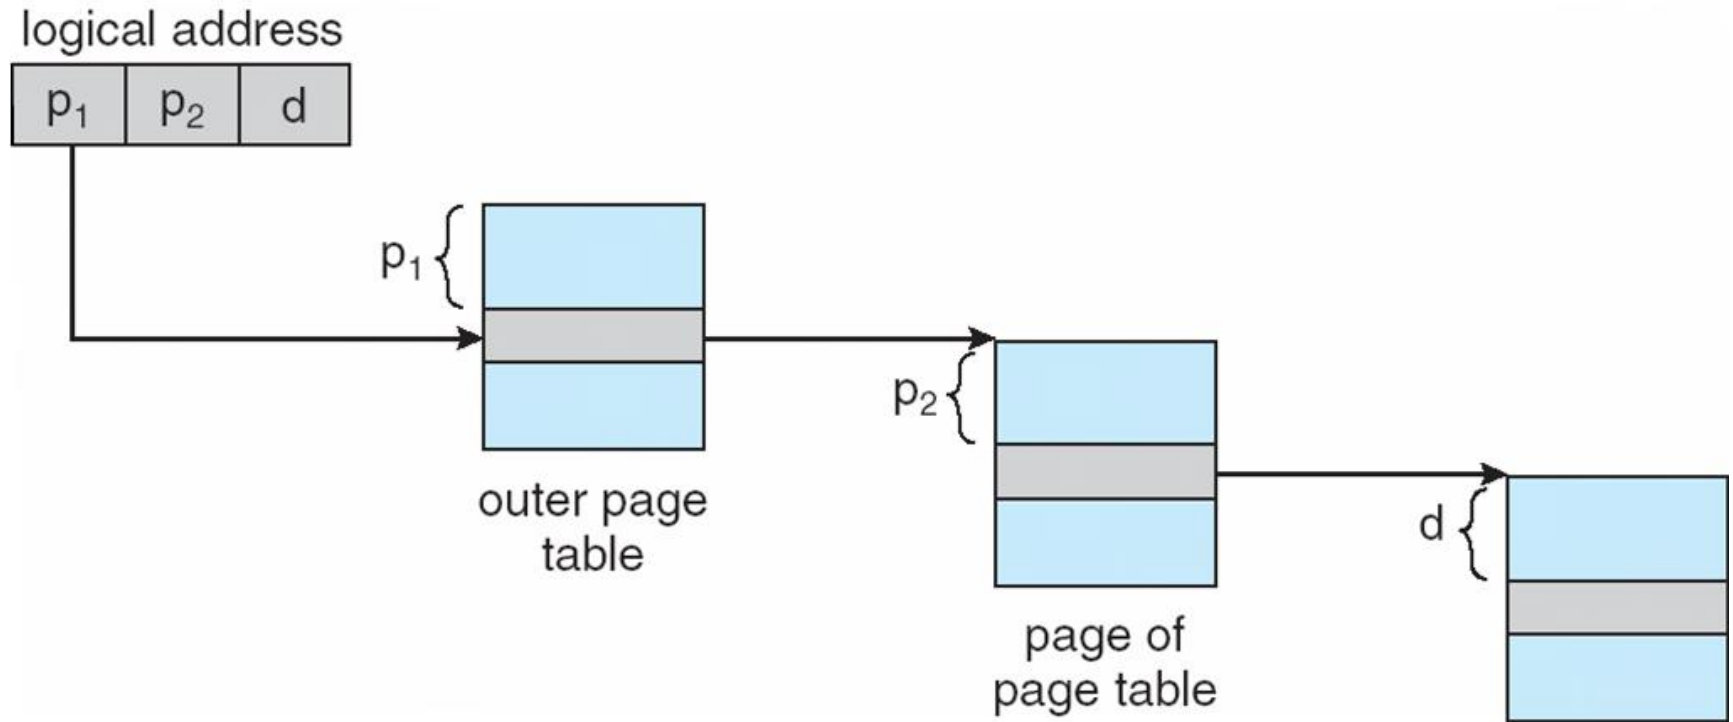
# Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
  - a page number consisting of 22 bits
  - a page offset consisting of 10 bits
- Since the page table is paged, then the page number is further divided into:
  - a 12-bit page number
  - a 10-bit page offset
- Thus, a logical address is as follows:
  - where  $p_1$  is an index into outer page table, and  $p_2$  is the displacement within the page of inner page table





# Address-Translation Scheme





# 64-bit Logical Address Space

- Even two-level paging scheme not sufficient
- If page size is 4 KB ( $2^{12}$ )
  - Then page table has  $2^{52}$  entries
  - If two level scheme, inner page tables could be  $2^{10}$  4-byte entries
  - Address would look like:

outer page	inner page	page offset
$p_1$	$p_2$	$d$
42	10	12
  - Outer page table has  $2^{42}$  entries
  - One solution is to add a 2<sup>nd</sup> outer page table
  - But in the following example the 2<sup>nd</sup> outer page table is still  $2^{32}$  bytes in size
    - ▶ And possibly 4 memory access to get to one physical memory location







# Three-level Paging Scheme

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12






# Hashed Page Tables

- Common in address spaces  $> 32$  bits
- The virtual page number is hashed into a page table
  - This page table contains a chain of elements hashing to the same location
- Each element contains:
  - (1) the virtual page number
  - (2) the value of the mapped page frame
  - (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
  - When a match is found, the corresponding physical frame is extracted






# Hashed Page Table





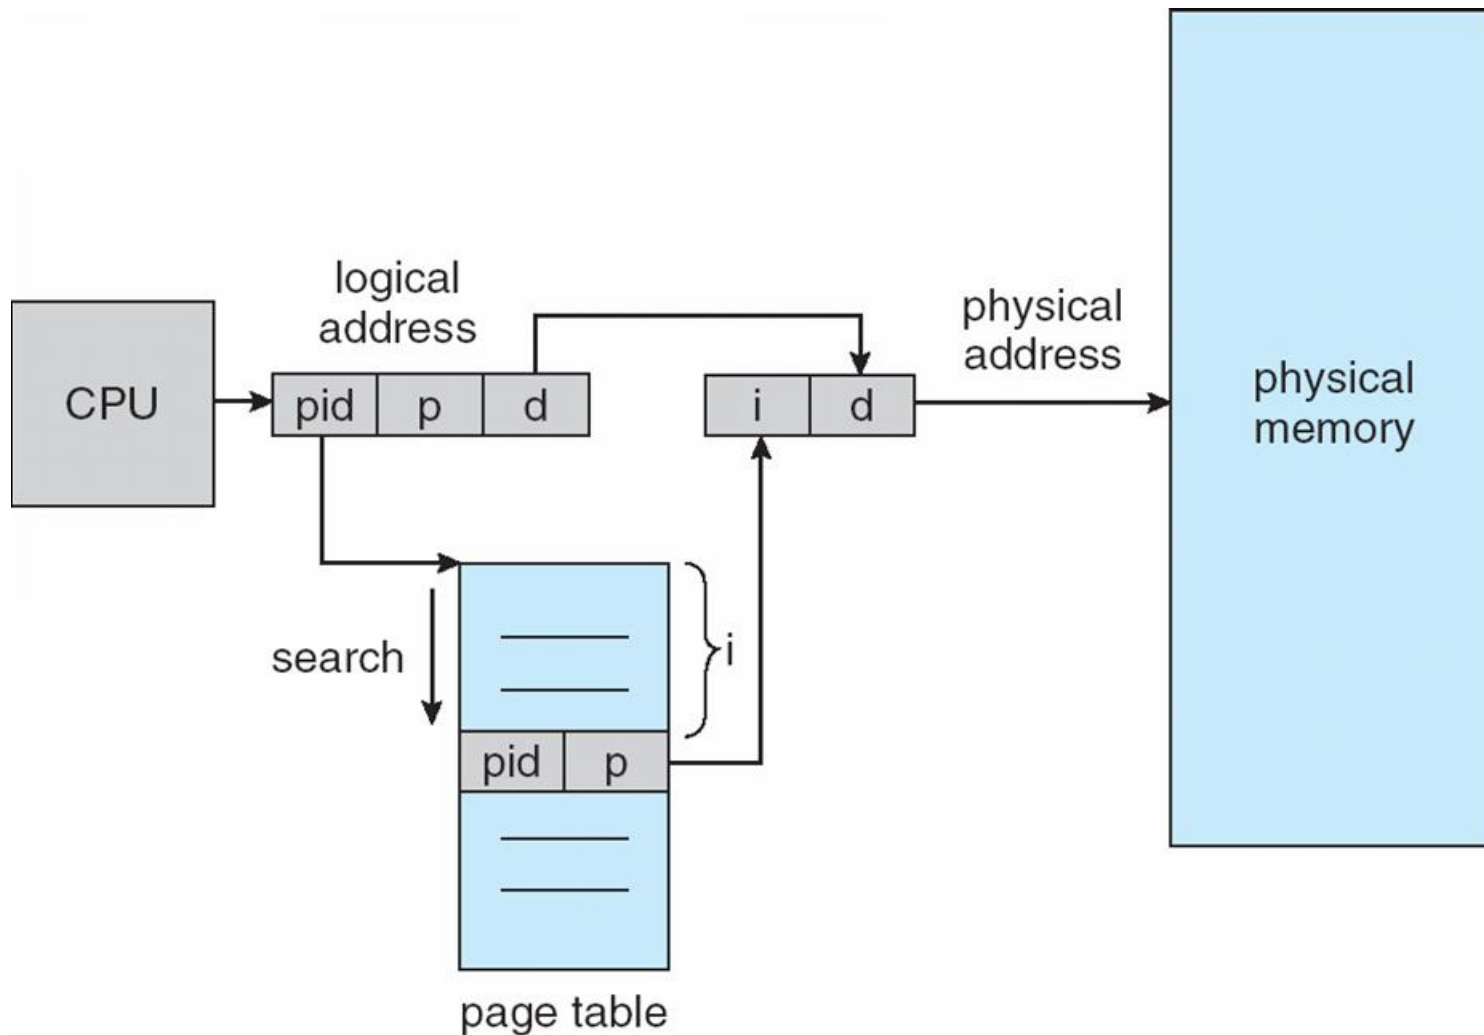
# Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
  - One entry for each real page (frame) of memory
  - Each entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
  - Use hash table to limit search to one, or at most a few, page-table entries and use TLB to accelerate access
- But how to implement shared memory?
  - One mapping of a virtual address to the shared physical address





# Inverted Page Table Architecture



# End of Chapter 8

---

